

ISO/IEC 42001:2023 AI Management System (AIMS) Mapping

Document Reference: DOC-02

System: Case Ace v2.0 Client Consultation Recording & Note Drafting System

Organisation: Citizens Advice Wandsworth (CAW)

Standard: ISO/IEC 42001:2023 (Information Technology — Artificial Intelligence — Management System)

Status: Aligned & Empirically Verified

Classification: Official-Sensitive / Governance Pack

1. Executive Statement & Scope Notice

[!IMPORTANT]

Declaration of AI Management Alignment:

This document details the technical, governance, and operational controls implemented in Case Ace v2.0 in accordance with ISO/IEC 42001:2023.

No claim of formal accredited certification is made. This mapping documents CAW's implementation of responsible, transparent, human-in-the-loop generative AI for advice consultation note-taking.

2. ISO/IEC 42001:2023 Control Domain Mapping

Domain 1: AI Policy & Organisational Objectives (Clause 5 & 6)

Objective: Establish organizational boundaries, ethical guardrails, and acceptable use policies for AI deployment.

CAW Policy Implementation:

Case Ace v2.0 is strictly constrained to administrative case note drafting from consented advice consultations.

The AI is explicitly forbidden from providing autonomous advice, determining client entitlement, making welfare benefit decisions, or contacting third parties.

Professional responsibility for the accuracy, completeness, and clinical/legal validity of the case note remains 100% with the qualified advising professional.

Evidence Reference: evidence/model_card.md, backend/src/prompts/caseRecordingMasterPrompt.ts.

Domain 2: AI Impact Assessment & Risk Governance (Clause 6.1.2 & Annex B)

Objective: Systematically identify, evaluate, and mitigate risks associated with AI deployment (e.g. hallucination, omission, demographic bias, automation bias, privacy breaches).

Identified Risks & Controls:

Automation Bias: Advisers blindly trusting AI-generated drafts.

Control: Deliberate UI friction in Phase 14 review interface. Mandatory individual confirmation of low-confidence statements, un-ticked gap acknowledgements, and explicit statutory responsibility confirmation before copy/export.

Harmful Hallucination / Omission: Fabricating debt figures, benefit deadlines, or advice given.

Control: Zero-temperature deterministic prompt structure, few-shot grounding in AQS Level 3 standards, explicit gap detection flagging missing documents/facts, and dual-pass verification.

Demographic & Dialect Bias: Poor ASR/NER performance on regional accents or non-native English speakers.

Control: 33-scenario synthetic corpus representing Geordie, Glaswegian, Cockney, Welsh, multicultural accents, and speech impairments. Redaction recall verified at $\geq 92.3\%$ and AQS compliance at 100%.

Data Exfiltration & Surveillance: Third-party LLM training on client disclosures.

Control: Local surrogate tokenisation before network egress. Zero-day retention agreement with Google Cloud Vertex AI in europe-west2 (London) under enterprise terms ensuring zero model training on client payloads.

Evidence Reference: evidence/residual_risk_register.md, evidence/redaction_performance_report.md.

Domain 3: AI Data Governance & Quality Management (Clause 8 & Annex B.4)

Objective: Ensure integrity, provenance, privacy, and suitability of data used across the AI lifecycle.

Data Lifecycle Controls:

Zero Production Client Training Data: Case Ace v2.0 utilizes frozen foundation models (gemini-1.5-pro-002 / gemini-1.5-flash-002) accessed strictly via inference APIs. No client data is ever used for model fine-tuning, training, or distillation.

Synthetic Test Corpus Governance: All 33 test scenarios are 100% synthetic, authored by CAW subject-matter experts with known ground truth spans and model case notes.

Local Tokenisation Boundary: Direct PII (NINOs, Names, Addresses, Phones, Employers) is replaced with surrogate tokens ([CLIENT_NAME_1], [NINO_1]) within the browser before the text reaches any AI model.

Evidence Reference: test/corpus/syntheticAdviceCorpus.ts, client/src/redaction/identifierEngine.ts.

Domain 4: Human Oversight, Agency & Professional Responsibility (Annex B.5)

Objective: Guarantee that human agents retain meaningful control, override capability, and ultimate accountability for AI outputs.

Phase 14 Adviser Review & Sign-Off Architecture:

Fully Editable Draft: Advisers can edit, delete, or append any section of the generated draft.

Bi-Directional Source Grounding: Clicking any sentence in the draft highlights the exact source utterance in the transcript; clicking a transcript segment displays the corresponding drafted text.

Mandatory Individual Acknowledgements: Gaps, limitations, and low-confidence statements are highlighted with required individual review actions. No "Accept All" or pre-ticked checkboxes exist.

Prominent Safeguarding Protocol: Safeguarding disclosures are presented in a dedicated high-visibility panel linking directly to CAW Safeguarding Procedures and designated safeguarding leads (DSLs).

Explicit Affirmation of Responsibility: Sign-off requires clicking a clear affirmative statement:

"I confirm that this case record is accurate, complete, and represents my professional record of this consultation. I acknowledge that professional responsibility for this advice rests solely with me."

Evidence Reference: client/src/components/AdviserReviewSignOff.tsx, evidence/model_card.md.

Domain 5: Transparency, Explainability & Auditability (Annex B.7 & B.8)

Objective: Ensure AI processing is transparent to clients and advisers, and operations are auditable without compromising privacy.

Transparency Mechanisms:

Tiered Client Consent (Phase 6): Clear, plain-English explanation provided to clients before consultation start, detailing audio processing, redaction, and right to withdraw consent at any time.

System Auditability: Immutable audit log store tracking prompt versions (v2.4.0), model identifiers (gemini-1.5-pro-002), verification pass results, and review gate durations without capturing free text or PII.

Evidence Reference: backend/src/logging/logSchema.ts, evidence/ropa_entry.md.

Domain 6: AI System Lifecycle Management & Continuous Validation (Clause 9 & 10)

Objective: Continuous monitoring, performance evaluation, and regression testing across the system lifecycle.

Continuous Testing Pipeline:

Automated benchmark execution on every commit evaluating 33 synthetic scenarios against AQS Level 3 quality criteria and redaction recall/precision thresholds (≥ 90%).

Fail-closed network egress interception asserting zero PII leakage.

Evidence Reference: test/testingEngine.ts, scripts/test-phase15-17.mjs.